

Machine learning based ranking for risk

Intelligent risk recommendations for ACS

Stephan Heßelmann

2025-11-05

Agenda

1. **The Challenge:** Why we need machine learning for risk ranking
2. **The Algorithm:** How a random forest works
3. **The Training:** From Central data to trained model
4. **The Predictions:** How to rank deployments in production
5. **Vibe Coding:** Current prototype status

THE CHALLENGE

What Is Risk?

- ACS monitors thousands of deployments across customer clusters
- Each deployment has dozens of security characteristics:
 - Policy violations
 - CVE vulnerabilities
 - Privileged containers
- ■ Network exposure
 - Process anomalies
 - ... and many more
- Question: Which deployments should security teams focus on first?

The Existing Risk Feature

Risk
Default view

Filter deployments

149 deployments

The ranking score is not explained at all.

Name	Created	Cluster	Namespace	Priority
! mid-server	28 Jan 2025, 0:09:33 CET	staging-secured-cluster	service-now	9
! mid-server	28 Jan 2025, 0:09:09 CET	staging-central-cluster	service-now	9
! rekor-redis	09 Apr 2025, 17:40:00 CEST	staging-central-cluster	trusted-artifact-signer	11
! visa-processor	21 Feb 2024, 22:07:13 CET	staging-secured-cluster	payments	12
! wordpress	28 Apr 2025, 17:36:32 CEST	staging-central-cluster	default	13
! nginx	09 Dec 2024, 23:54:32 CET	staging-central-cluster	default	17
✓ search-postgres	14 Oct 2025, 17:54:22 CEST	staging-central-cluster	open-cluster-management	39
✓ trillian-db	09 Apr 2025, 17:38:33 CEST	staging-central-cluster	trusted-artifact-signer	42

Current risk assessment in ACS

Shortcomings of Current Approach

-  **Intransparent**
 - The risk score is baked into ACS with a complex formula.
 - Customers cannot understand or validate how scores are calculated.
-  **Inflexible**
 - Does not adapt to customer needs or emerging trends.
 - One-size-fits-all approach across all environments.

Current Approach: Rule-Based Multipliers

Feature multiplication:

Risk Score =

Policy Violations ×
Process Baseline ×
Vulnerabilities ×
Risky Components ×
Component Count ×
Image Age ×
...

Drawbacks:

- **Hardcoded formulas**
 - Manually tuned thresholds and weights
- **Cannot reverse the calculation**
 - Multiplied scores lose individual feature information

Example: Image Age Calculation

```
1 // Creates a score that is:  
2 // A) No risk when daysSinceCreated is < penalizeDaysFloor  
3 if imageAgeInDays < penalizeDaysFloor {  
4     return  
5 }  
6  
7 // B) Increases linearly between penalizeDaysFloor and penalizeDaysCeil  
8 daysSincePenalized := imageAgeInDays - penalizeDaysFloor  
9 scaledDays := float32(daysSincePenalized) / float32(penalizeDaysCeil-penali  
10 riskScore = float32(1) + float32(scaledDays*(maxMultiplier-1))  
11  
12 // C) Is 1.5 when duration is > penalizeDaysCeil  
13 if riskScore > maxMultiplier {  
14     riskScore = maxMultiplier  
15 }
```

The Machine Learning Solution

- **Learn** from real or synthetic training data.
- **Identify** which security features matter most.
- **Rank** deployments by learned feature weights.
- **Explain** what features contribute to high risk scores.
- **Improve** continuously as new data arrives.


THE ALGORITHM

Random Forest: The Committee Approach

Analogy: Assessing fire risks for a house.

- Instead of **1 inspector** with a rigid checklist...
 - ... ask **1000 inspectors**.
- Each inspector looks at different factors:
 - Electrical wiring + smoke detectors
 - Flammable materials + escape routes
 - Building age + maintenance history
- **Final assessment** = average of all 1000 opinions

Random Forest Architecture



Example: One Decision Tree

```
Is vulnerability_score > 7.5?  
└ YES → Is privileged_container_ratio > 0.5?  
    └ YES → Risk Score = 8.2  
    └ NO → Is external_exposure = 1?  
        └ YES → Risk Score = 6.5  
        └ NO → Risk Score = 4.1  
└ NO → Is policyViolation_score > 3.0?  
    └ YES → Risk Score = 5.0  
    └ NO → Risk Score = 2.3
```

Final Prediction: Average of 1000s such trees

Why Random Forest?

1.  **Robust against outliers**
 - If one tree makes a mistake, 999 others can correct it.
2.  **Handles complexity**
 - Learns non-linear relationships automatically.
3.  **Provides Explanations**
 - Feature importance: which security issues matter most.
4.  **Simple Training**
 - No epochs, no backpropagation, no GPUs required.

Why not Random Forest?

1. ~~X~~ Just a regression algorithm

- No actual AI. Not in the training data -> won't be learned.

2. ~~X~~ Supervised learning


- Predictions are only as good as the training input. Feature engineering matters.

3. ~~X~~ Retraining

- Learning new things requires new training data + model retraining.


THE TRAINING

End-to-End Training Flow



THE PREDICTIONS

Prediction Flow



Feature Importance: What Matters?

Example: Top 5 most important features

1. policyViolationScore	28.5%
2. maxVulnerabilityScore	19.8%
3. privilegedContainerRatio	15.2%
4. externalExposure	8.9%
5. processBaselineViolations	7.6%

- **Insight:** Policy violations and critical CVEs drive risk ranking more than other factors.
- **Detail:** Calculated for training data and each prediction.

Performance Metric: NDCG

NDCG = Normalized Discounted Cumulative Gain

- Measures **ranking quality** on 0-1 scale
 - **1.0**: perfect ranking, **0.0**: worst ranking

Example:

True ranking: [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

Model ranking: [1, 2, 3, 4, 5, 6, 7, 9, 8, 10]

↑ Small error at bottom

NDCG ≈ 0.98 (98% as good as perfect ranking)

Typical Performance

Training Metrics

- NDCG: **0.75-0.95**
- Training time: **~60 sec**

Model Characteristics

- Feature count: **~20**
- Training samples: **100-1000**
- Model size: **~20 MB**
- Trees: **1000**

Validation **NDCG > 0.75** means model ranks deployments reasonably well!

VIBE CODING

Current Prototype

Model

- Training pipeline with Central integration
- Feature extraction for deployments + images
- RandomForest model with versioning
- Validation with NDCG tracking

Deployment

- Python container with FastAPI REST service
- Model persistence to local/remote storage

Questions?

Future Improvements

- Fine-tune different models and improve feature extraction.
- Use LLMs to generate advanced training data.
- Implement back propagation of user driven risk decisions.
- Integrate prototype with ACS UI.

Resources

- Code: [stackrox/ml-risk-service](#)
- Docs: [Learning.md](#)